

MACHINE LEARNING

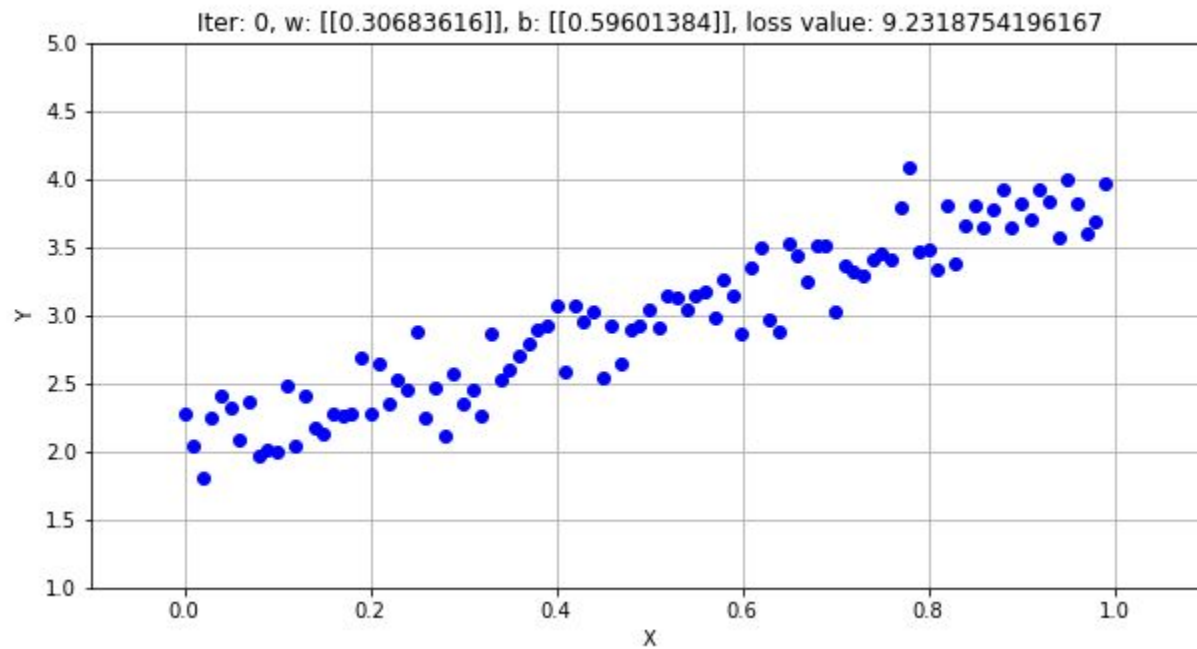
LINEAR REGRESSION

Professor Ernesto Lee



[HTTPS://BIT.LY/MLTRAIN](https://bit.ly/mltrain)

Go here for your LAB ENVIRONMENT

LINEAR RELATIONSHIPS



REGRESSION ALGORITHM



LINEAR REGRESSION WITH SCIKIT-LEARN

IMPORT YOUR LIBRARIES

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split, cross_val_score
```

```
from sklearn import metrics
```


PULL IN YOUR DATA

You can access your data here:

<https://bit.ly/21homes>

Or here:

<https://raw.githubusercontent.com/fenago/pythonml/main/data/HousePrice.csv>

```
df =  
pd.read_csv('https://raw.githubusercontent.com/fenago/pythonml/  
main/data/HousePrice.csv')
```

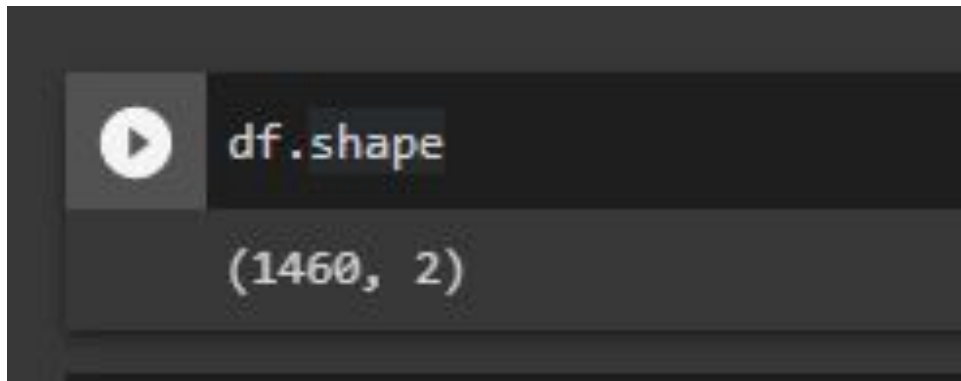
PUTTING THE COLUMNS YOU WANT IN YOUR DATA

```
df2 = df[["open", "close"]] #open and close are the  
columns
```

Make sure you have the columns you want and need!

UNDERSTAND YOUR DATA WITH SHAPE

df.shape

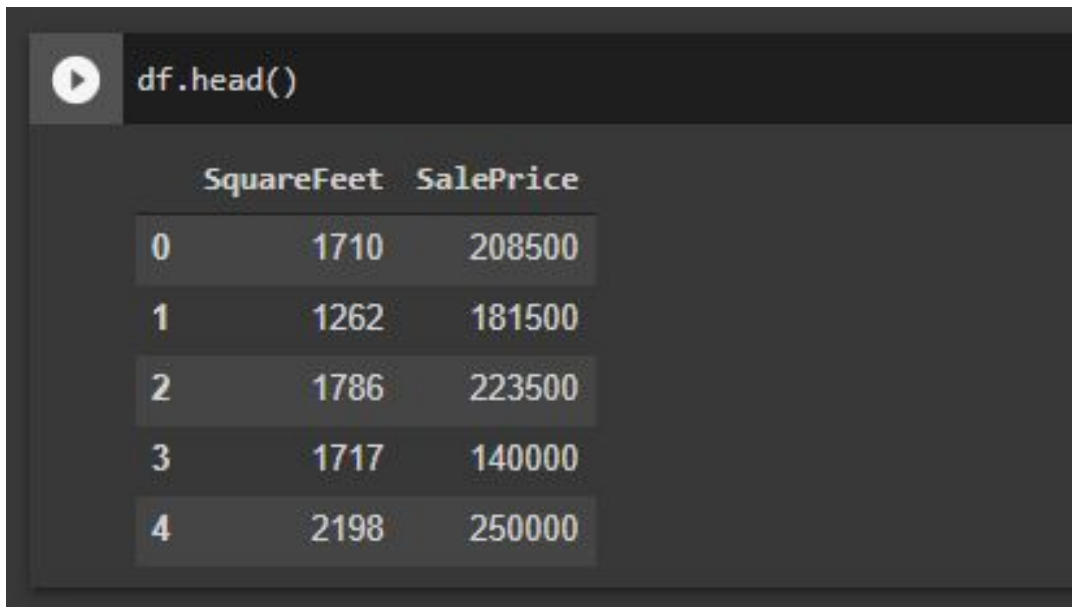
A screenshot of a Jupyter Notebook cell. The cell contains the code `df.shape` and its output, `(1460, 2)`. The code is highlighted with a light blue background, and the output is displayed below it. A play button icon is visible on the left side of the cell.

```
df.shape
```

```
(1460, 2)
```

UNDERSTAND YOUR DATA: HEAD

df.head()

A screenshot of a Jupyter Notebook cell. The cell contains the code `df.head()` and its output. The output is a table with 5 rows and 3 columns. The first column is an index from 0 to 4. The second column is labeled 'SquareFeet' and the third is labeled 'SalePrice'.

	SquareFeet	SalePrice
0	1710	208500
1	1262	181500
2	1786	223500
3	1717	140000
4	2198	250000

UNDERSTAND YOUR DATA: DESCRIPTIVE STATISTICS

`df.describe()`

```
[8] df.describe()
```

	SquareFeet	SalePrice
count	1460.000000	1460.000000
mean	1515.463699	180921.195890
std	525.480383	79442.502883
min	334.000000	34900.000000
25%	1129.500000	129975.000000
50%	1464.000000	163000.000000
75%	1776.750000	214000.000000
max	5642.000000	755000.000000


FIND YOUR CORRELATIONS IN YOUR DATASETS

correlation between 2 Specific Columns


```
print(df['SquareFeet'].corr(df['SalesPrice']))
```

pair-wise correlation between all columns


```
print(df.corr())
```




Strong Positive Correlation




Weak Positive Correlation



No Correlation



Strong Negative Correlation



Weak Negative Correlation

HEATMAP

```
# Correlation between different variables
corr = df.corr()

# Set up the matplotlib plot configuration
f, ax = plt.subplots(figsize=(12, 10))

# Generate a mask for upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Configure a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap
sns.heatmap(corr, annot=True, mask = mask, cmap=cmap)
```


VISUALIZE YOUR DATA: PLOT

```
df.plot(x='SquareFeet', y='SalePrice', style='*')  
plt.title('Square Feet vs Sale Price')  
plt.xlabel('Square Feet')  
plt.ylabel('Sale Price')  
plt.show()
```

PREPARE YOUR DATA: SPLIT INTO TRAINING AND TEST SETS

```
X = df.iloc[:, :-1].values
```

```
y = df.iloc[:, 1].values
```

TRAIN TEST SPLIT

```
X_train, X_test, y_train, y_test = train_test_split(X,  
y, test_size=0.2, random_state=0)
```

RUN THE MODEL

```
def get_cv_scores(model):  
    scores = cross_val_score(model, X_train, y_train, cv=10, scoring='r2')  
  
    print('CV Mean: ', np.mean(scores))  
    print('STD: ', np.std(scores))  
    print('\n')  
  
lr = LinearRegression().fit(X_train, y_train)  
get_cv_scores(lr)
```

VIEW THE COEFFICIENTS

```
print(lr.intercept_)
```

```
print(lr.coef_)
```



```
print(lr.intercept_)  
print(lr.coef_)
```

```
13330.293444921088  
[110.26434426]
```


PREDICTIONS

```
y_pred = lr.predict(X_test)
```

```
plt.scatter(X_train, y_train)
```

```
plt.plot(X_test, y_pred, color='red')
```

```
plt.show()
```



PREDICT WITH NEW UNSEEN DATA

```
lr.predict([[2515]])
```

EVALUATE PREDICTED VALUES FROM ACTUALS

```
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
```

```
df.head()
```



```
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})  
df.head()
```

	Actual	Predicted
0	200624	290645.119259
1	133000	187327.428687
2	110000	145978.299590
3	192000	236284.797539
4	88000	133738.957377

R-SQUARED (R²)

```
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})  
df.head()  
  
print('R-Squared:', metrics.r2_score(df['Actual'], df['Predicted']))  
-----
```

Summary Definition

Define R-Squared: Coefficient of determination means a statistical measurement of the correlation between two variables.

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.r2_score.html

EVALUATE YOUR MODEL

```
print('Mean Absolute Error:',  
metrics.mean_absolute_error(y_test, y_pred))
```

```
print('Mean Squared Error:',  
metrics.mean_squared_error(y_test, y_pred))
```

```
print('Root Mean Squared Error:',  
np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```



```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
Mean Absolute Error: 39364.76724953735
```

```
Mean Squared Error: 3913788296.4027987
```

```
Root Mean Squared Error: 62560.277304394986
```

<https://scikit-learn.org/>

<https://github.com/scikit-learn/scikit-learn>

TUNING HYPERPARAMETERS WITH SIMPLE LINEAR REGRESSION

<https://bit.ly/ucidata>

Hint: To create a new dataframe with selected columns - do this:

```
df2 = df[["open", "close"]] #open and close are the  
columns
```

LAB 1

Read in this DATA:

```
pd.read_csv('https://raw.githubusercontent.com/fenago/pythonml/main/data/poverty.txt', sep="\t")
```

Apply what you have learned to create Simple Linear Models.

This dataset of size $n = 51$ are for the 50 states and the District of Columbia in the United States ([poverty.txt](#)). The variables are y = year 2002 birth rate per 1000 females 15 to 17 years old and x = poverty rate, which is the percent of the state's population living in households with incomes below the federally defined poverty level. (Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard).

LAB 2

Read in this DATA:

```
pd.read_csv('https://raw.githubusercontent.com/fenago/pythonml/main/data/lungfunction.txt', sep="\t")
```

Apply what you have learned to create Simple Linear Models.

This dataset of size $n = 51$ are for the 50 states and the District of Columbia in the United States ([poverty.txt](#)). The variables are y = year 2002 birth rate per 1000 females 15 to 17 years old and x = poverty rate, which is the percent of the state's population living in households with incomes below the federally defined poverty level. (Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard).

MULTIVARIATE LINEAR REGRESSION

OUR DATASET

Data Dictionary:

<http://people.sc.fsu.edu/~jburkardt/datasets/regression/x16.txt>

Dataset:

```
pd.read_csv("https://raw.githubusercontent.com/fenago/pythonml/main/data/petrol\_consumption.csv")
```

LOAD THE LIBRARIES

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

LOAD THE DATA

```
df =  
pd.read_csv("https://raw.githubusercontent.com/fenago/python-ml/main/data/petrol\_consumption.csv")  
  
df.head()  
  
df.describe()
```

FIND YOUR CORRELATIONS IN YOUR DATASETS

```
# correlation between 2 Specific Columns
```

```
print(df['Petrol_tax'].corr(df['Petrol_Consumption']))
```

```
# pair-wise correlation between all columns
```

```
print(df.corr())
```

PREPARE THE DATA

```
X = dataset[['Petrol_tax', 'Average_income',  
            'Paved_Highways', 'Population_Driver_licence(%)']]
```

```
y = dataset['Petrol_Consumption']
```

```
#Execute below to divide into train/test sets
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.2, random_state=0)
```

TRAIN THE ALGORITHM AS BEFORE

```
from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

WHAT COEFFICIENTS DID IT FIND?

```
coeff_df = pd.DataFrame(regressor.coef_, X.columns,  
columns=['Coefficient'])
```

coeff_df

	Coefficient
Petrol_tax	-24.196784
Average_income	-0.81680
Paved_Highways	-0.000522
Population_Driver_license(%)	1324.675464

PREDICTIONS

```
y_pred = regressor.predict(X_test)
```

```
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
```

df

	Actual	Predicted
36	640	643.176639
22	464	411.950913
20	649	683.712762
38	648	728.049522
18	865	755.473801

EVALUATE THE ALGORITHM

```
from sklearn import metrics

print('Mean Absolute Error:',
metrics.mean_absolute_error(y_test, y_pred))

print('Mean Squared Error:',
metrics.mean_squared_error(y_test, y_pred))


print('Root Mean Squared Error:',
np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

print('R-Squared:', metrics.r2_score(df['Actual'],
df['Predicted']))
```


LOGISTIC REGRESSION

LOGISTIC REGRESSION

LOGISTIC REGRESSION




LOGISTIC REGRESSION

when $b_0 + b_1x = 0$, then
the p will be 0.5,

similarly, $b_0 + b_1x > 0$,
then the p will be going
towards 1 and

$b_0 + b_1x < 0$, then the p
will be going towards 0.



LOADING THE DATA


```
from sklearn.datasets import load_digits  
  
digits = load_digits()
```

SHOWING THE IMAGES AND LABELS

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(20,4))
```



```
for index, (image, label) in enumerate(zip(digits.data[0:5], digits.target[0:5])):
```

```
    plt.subplot(1, 5, index + 1)
```

```
    plt.imshow(np.reshape(image, (8,8)), cmap=plt.cm.gray)
```

```
    plt.title('Training: %i\n' % label, fontsize = 20)
```


SPLIT THE DATA INTO TRAINING AND TEST SETS

```
from sklearn.model_selection import train_test_split  
  
x_train, x_test, y_train, y_test =  
train_test_split(digits.data, digits.target, test_size=0.25,  
random_state=0)
```

MODELING

```
from sklearn.linear_model import LogisticRegression  
  
logisticRegr = LogisticRegression()  
  
logisticRegr.fit(x_train, y_train)
```

```
#predict for one image
```

```
logisticRegr.predict(x_test[0].reshape(1,-1))
```

```
#predict for multiple images
```

```
logisticRegr.predict(x_test[0:10])
```

```
#for the entire dataset
```

```
predictions = logisticRegr.predict(x_test)
```

```
[14] #predict for one image
      logisticRegr.predict(x_test[0].reshape(1,-1))
```

```
array([2])
```

```
[15] #predict for multiple images
      logisticRegr.predict(x_test[0:10])
```

```
array([2, 8, 2, 6, 6, 7, 1, 9, 8, 5])
```

```
[16] #for the entire dataset
      predictions = logisticRegr.predict(x_test)
```

```
[17] predictions
```

```
array([2, 8, 2, 6, 6, 7, 1, 9, 8, 5, 2, 8, 6, 6, 6, 6, 1, 0, 5, 8, 8, 7,
       8, 4, 7, 5, 4, 9, 2, 9, 4, 7, 6, 8, 9, 4, 3, 1, 0, 1, 8, 6, 7, 7,
       1, 0, 7, 6, 2, 1, 9, 6, 7, 9, 0, 0, 9, 1, 6, 3, 0, 2, 3, 4, 1, 9,
       2, 6, 9, 1, 8, 3, 5, 1, 2, 8, 2, 2, 9, 7, 2, 3, 6, 0, 5, 3, 7, 5,
       1, 2, 9, 9, 3, 1, 4, 7, 4, 8, 5, 8, 5, 5, 2, 5, 9, 0, 7, 1, 4, 7,
       3, 4, 8, 9, 7, 9, 8, 2, 1, 5, 2, 5, 8, 4, 1, 7, 0, 6, 1, 5, 5, 9,
       9, 5, 9, 9, 5, 7, 5, 6, 2, 8, 6, 9, 6, 1, 5, 1, 5, 9, 9, 1, 5, 3,
       6, 1, 8, 9, 8, 7, 6, 7, 6, 5, 6, 0, 8, 8, 9, 9, 6, 1, 0, 4, 1, 6,
       3, 8, 6, 7, 4, 9, 6, 3, 0, 3, 3, 3, 0, 7, 7, 5, 7, 8, 0, 7, 1, 9,
       6, 4, 5, 0, 1, 4, 6, 4, 3, 3, 0, 9, 5, 9, 2, 8, 4, 2, 1, 6, 8, 9,
       2, 4, 9, 3, 7, 6, 2, 3, 3, 1, 6, 9, 3, 6, 3, 3, 2, 0, 7, 6, 1, 1,
       9, 7, 2, 7, 8, 5, 5, 7, 5, 3, 3, 7, 2, 7, 5, 5, 7, 0, 9, 1, 6, 5,
       9, 7, 4, 3, 8, 0, 3, 6, 4, 6, 3, 2, 6, 8, 8, 8, 4, 6, 7, 5, 2, 4,
       5, 3, 2, 4, 6, 9, 4, 5, 4, 3, 4, 6, 2, 9, 0, 1, 7, 2, 0, 9, 6, 0,
       4, 2, 0, 7, 9, 8, 5, 7, 8, 2, 8, 4, 3, 7, 2, 6, 9, 9, 5, 1, 0, 8,
       2, 8, 9, 5, 6, 2, 2, 7, 2, 1, 5, 1, 6, 4, 5, 0, 9, 4, 1, 1, 7, 0,
       8, 9, 0, 5, 4, 3, 8, 8, 6, 5, 3, 4, 4, 4, 8, 8, 7, 0, 9, 6, 3, 5,
       2, 3, 0, 8, 8, 3, 1, 3, 3, 0, 0, 4, 6, 0, 7, 7, 6, 2, 0, 4, 4, 2,
       3, 7, 1, 9, 8, 6, 8, 5, 6, 2, 2, 3, 1, 7, 7, 8, 0, 3, 3, 1, 1, 5,
       5, 9, 1, 3, 7, 0, 0, 3, 0, 4, 5, 8, 9, 3, 4, 3, 1, 8, 9, 8, 3, 6,
       3, 1, 6, 2, 1, 7, 5, 5, 1, 9])
```

EVALUATION

#Accuracy = correct predictions / total number of data points

#Use score method to get accuracy of model

```
score = logisticRegr.score(x_test, y_test)
```

```
print(score)
```

```
[18] #Use score method to get accuracy of model  
      score = logisticRegr.score(x_test, y_test)  
      print(score)
```

```
0.9511111111111111
```

CONFUSION MATRIX

```
from sklearn import metrics
```

```
cm = metrics.confusion_matrix(y_test, predictions)
```

```
print(cm)
```

```
[19] from sklearn import metrics
```


```
cm = metrics.confusion_matrix(y_test, predictions)  
print(cm)
```

```
[[37  0  0  0  0  0  0  0  0  0]  
 [ 0 40  0  0  0  0  0  0  2  1]  
 [ 0  1 40  3  0  0  0  0  0  0]  
 [ 0  0  0 43  0  0  0  0  1  1]  
 [ 0  0  0  0 37  0  0  1  0  0]  
 [ 0  0  0  0  0 46  0  0  0  2]  
 [ 0  1  0  0  0  0 51  0  0  0]  
 [ 0  0  0  1  1  0  0 46  0  0]  
 [ 0  3  1  0  0  0  0  0 43  1]  
 [ 0  0  0  0  0  1  0  0  1 45]]
```



KNN

KNN

Choice of value of K




HOW DOES KNN WORK?



DISTANCE IN KNN

$$\begin{aligned}d(\mathbf{p}, \mathbf{q}) &= d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2} \\&= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}.\end{aligned}$$

NEAREST?



STEPS TO SOLVE A KNN PROBLEM

1. Load and store the data.
2. Calculate the distance from x (new data point) to all other data points.
3. Sort all the distances from your data in ascending order.
4. Initialize the K value for the nearest data points.
5. Make a prediction based on the majority of data points with the same label within the K value.
6. Evaluate your machine learning model.

THE USE CASE...

we want to create a machine model that will allow botanists to classify different species of iris flowers.



LET'S BUILD THE MODEL

IMPORTS

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib
```

DATA PREPARATION

```
from sklearn.datasets import load_iris  
  
iris_dataset = load_iris()  
  
print("Keys of iris_dataset:\n", iris_dataset.keys())
```

OUTPUT

```
Keys of iris_dataset:  
  
dict_keys(['data', 'target', 'target_names', 'DESCR',  
'feature_names', 'filename'])
```

DATA TYPE

```
print("Data Type:", type(iris_dataset['data']))
```

OUTPUT:

```
Data Type: <class 'numpy.ndarray'>
```

SHAPE

```
print("Shape of Data:", iris_dataset['data'].shape)
```

FEATURES

```
print("First 10 Samples and Their Features:\n",  
iris_dataset['data'][:10])
```

TARGET

```
print("Type of Target:", type(iris_dataset['target']))  
print("Shape of Target:", iris_dataset['target'].shape)  
print(iris_dataset['target'])
```

TARGET NAMES

```
print("Target names:", iris_dataset['target_names'])
```

DESCRIPTION

```
print(iris_dataset['DESCR'])
```

OR

```
print(iris_dataset['DESCR'][:500] + "\n...")
```


FEATURE NAMES

```
print("Feature Names:", iris_dataset['feature_names'])
```

TRAINING AND TESTING DATA

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(  
    iris_dataset['data'], iris_dataset['target'],  
    random_state=0)
```

VALIDATE THE SHAPE OF THE DATA AND TEST THE DATASET

```
print("X_train Shape:", X_train.shape)
```

```
print("y_train Shape:", y_train.shape)
```

```
print("X_test Shape:", X_test.shape)
```

```
print("y_test Shape:", y_test.shape)
```

VISUALIZE YOUR DATA

CREATE THE DF FOR VISUALIZING

```
df = pd.DataFrame(X_train,  
columns=iris_dataset.feature_names)
```

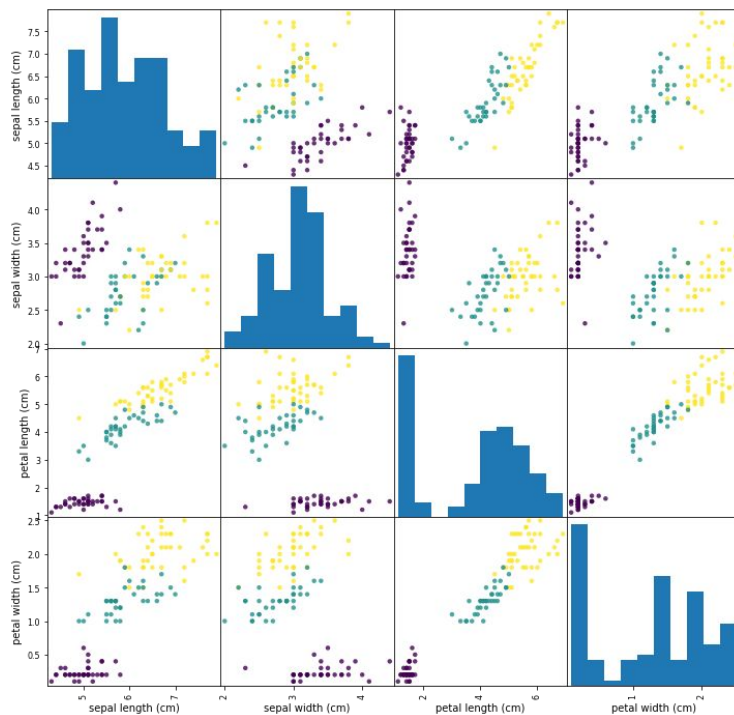
```
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.9	3.0	4.2	1.5
1	5.8	2.6	4.0	1.2
2	6.8	3.0	5.5	2.1
3	4.7	3.2	1.3	0.2
4	6.9	3.1	5.1	2.3

SHOW THE VISUALIZATION

```
pd.plotting.scatter_matrix(df,c=y_train,figsize=(12,12),  
marker='o',s=20,alpha=.8)
```

```
plt.show()
```

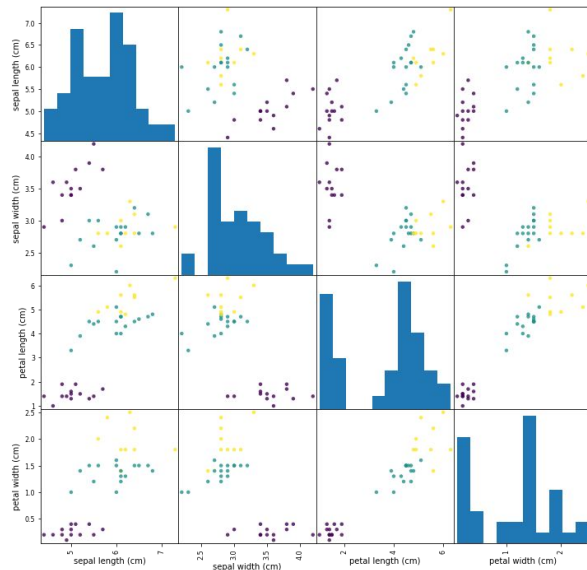


DO THE SAME WITH X

```
df = pd.DataFrame(X_test,  
columns=iris_dataset.feature_names)
```

```
pd.plotting.scatter_matrix(df,c=y_test,figsize=(12,12),  
marker='o',s=20,alpha=.8)
```

```
plt.show()
```



CREATE THE MODEL

```
from sklearn.neighbors import KNeighborsClassifier  
knnObject = KNeighborsClassifier(n_neighbors=1)  
knnObject.fit(X_train, y_train)
```

MAKING PREDICTIONS

PREDICTIONS ON NEW DATA

Imagine you are a machine learning engineer for a company. A client of yours reached out to you to verify an iris species they found in the wild. They only gave us the following information:

- Sepal length: 40 cm
- Sepal width: 10 cm
- Petal length: 5 cm
- Petal width: 2 cm

```
newIris = np.array([[40, 10, 5, 2]])  
print("newIris Shape:", newIris.shape)
```

PREDICT

```
prediction = knnObject.predict(newIris)
print("Prediction Value:", prediction)
print("Predicted Target Name:",
      iris_dataset['target_names'][prediction])
```

MODEL EVALUATION

SEE YOUR TEST PREDICTIONS

```
testSetPredictions = knnObject.predict(X_test)
print("Test Set Predictions:", testSetPredictions)
```

CHECK FOR ACCURACY

```
accuracy = round(knnObject.score(X_test, y_test),2)  
print("The Test Set Accuracy is:",accuracy)
```


BREAST CANCER USE CASE - KNN

INITIALIZE THE LIBRARIES

```
import pandas as pd
```

```
import numpy as np
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

IMPORT THE DATA SET

It is in the data folder and named dataR2.csv

```
data=pd.read_csv("./data/dataR2.csv")
```

UNDERSTAND YOUR DATA

`data.shape`

	Age	BMI	Glucose	Insulin	HOMA	Leptin	Adiponectin	Resistin	MCP.1	Classification
0	48	23.500000	70	2.707	0.467409	8.8071	9.702400	7.99585	417.114	1
1	83	20.690495	92	3.115	0.706897	8.8438	5.429285	4.06405	468.786	1
2	82	23.124670	91	4.498	1.009651	17.9393	22.432040	9.27715	554.697	1
3	68	21.367521	77	3.226	0.612725	9.8827	7.169560	12.76600	928.220	1
4	86	21.111111	92	3.549	0.805386	6.6994	4.819240	10.57635	773.920	1
...
111	45	26.850000	92	3.330	0.755688	54.6800	12.100000	10.96000	268.230	2
112	62	26.840000	100	4.530	1.117400	12.4500	21.420000	7.32000	330.160	2
113	65	32.050000	97	5.730	1.370998	61.4800	22.540000	10.33000	314.050	2
114	72	25.590000	82	2.820	0.570392	24.9600	33.750000	3.27000	392.460	2
115	86	27.180000	138	19.910	6.777364	90.2800	14.110000	4.35000	90.090	2

116 rows × 10 columns

FIND MISSING VALUES

```
data.isna().sum()
```

```
Age      0
BMI      0
Glucose  0
Insulin  0
HOMA     0
Leptin   0
Adiponectin  0
Resistin 0
MCP.1    0
Classification  0
dtype: int64
```

EXPLORATORY DATA ANALYSIS

data.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 116 entries, 0 to 115
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   116 non-null   int64
1   BMI                   116 non-null   float64
2   Glucose               116 non-null   int64
3   Insulin               116 non-null   float64
4   HOMA                  116 non-null   float64
5   Leptin                116 non-null   float64
6   Adiponectin           116 non-null   float64
7   Resistin              116 non-null   float64
8   MCP.1                 116 non-null   float64
9   Classification         116 non-null   int64
dtypes: float64(7), int64(3)
memory usage: 9.2 KB
```

HEATMAPS AND CORRELATION

#Heatmap to find correlation

```
plt.subplots(figsize=(20,20))
```

```
sns.heatmap(data.corr(), cmap='RdYlGn', annot=True)
```



COLUMNS


`data.columns`

KNN IS SENSITIVE TO
OUTLIERS

OUTLIERS FOR AGE?

#No outliers for age


```
sns.boxplot(data['Age'])
```



OUTLIERS FOR BMI?

```
#NO outliers for BMI
```

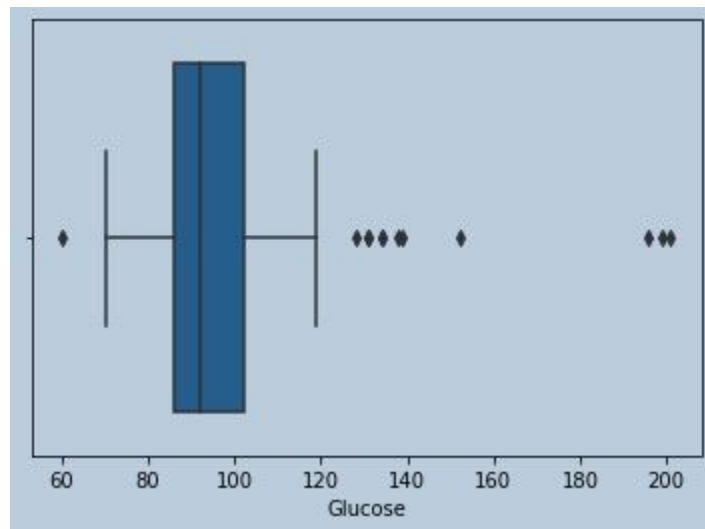
```
sns.boxplot(data['BMI'])
```



OUTLIERS FOR GLUCOSE

#Some outliers are there for Glucose and data is Skewed

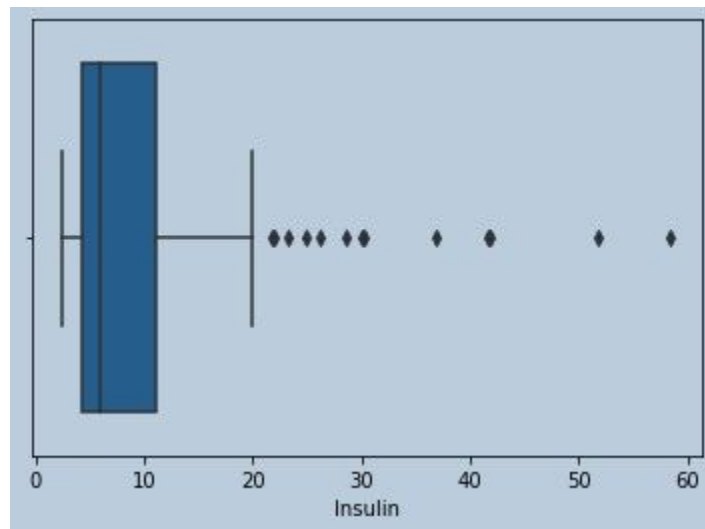
```
sns.boxplot(data['Glucose'])
```



OUTLIERS FOR INSULIN?

#Outliers are present in Insulin


```
sns.boxplot(data['Insulin'])
```



OUTLIERS FOR HOMA

#lots of Outliers in Homa

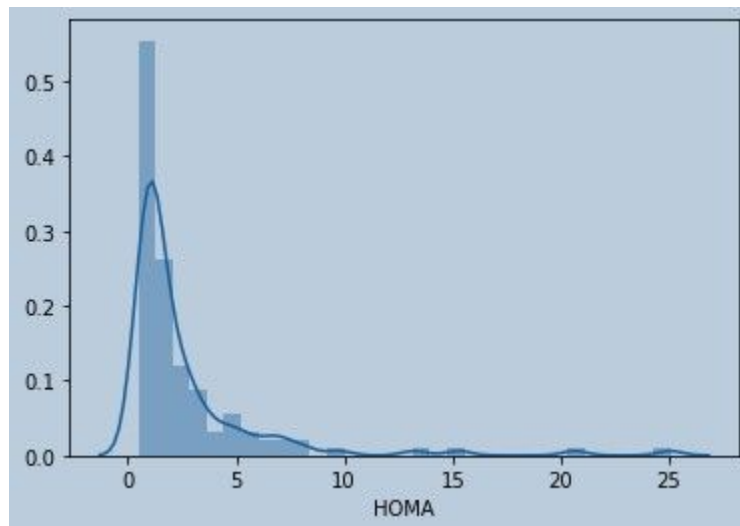
```
sns.boxplot(data['HOMA'])
```



DISTRIBUTION OF HOMA

#Distribution plot of HOMA

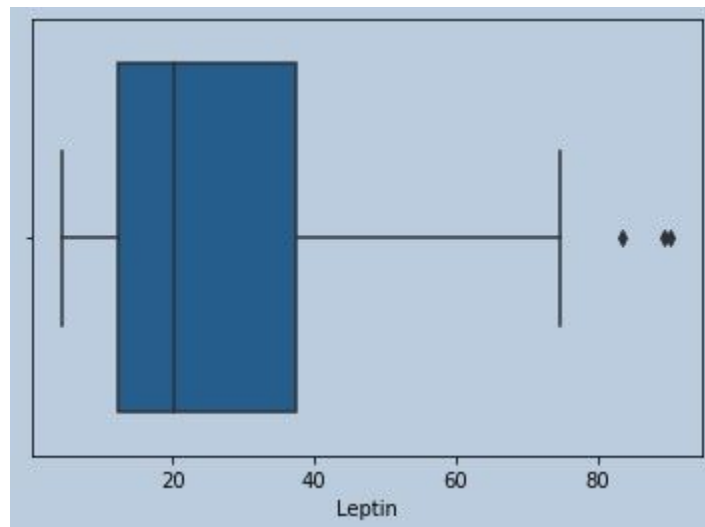
```
sns.distplot(data['HOMA'])
```



OUTLIERS FOR LEPTIN

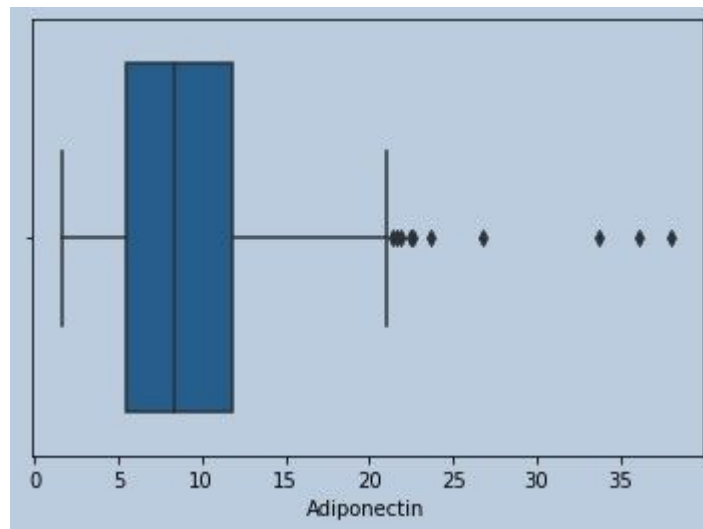
```
#Outliers present for Leptin
```

```
sns.boxplot(data['Leptin'])
```



#Outliers present for Adiponectin


```
sns.boxplot(data['Adiponectin'])
```



OUTLIERS FOR RESISTIN

#Outliers present for Resistin


```
sns.boxplot(data['Resistin'])
```



OUTLIERS FOR MCP.1?

```
#Outliers present for MCP.1
```

```
sns.boxplot(data['MCP.1'])
```



REMOVE OUTLIERS

#Removing Outliers Since they may affect prediction for KNN (quantile method)

```
cancer=data.copy()
```

```
insulinQ1=cancer['Insulin'].quantile(0.25)
```

```
insulinQ3=cancer['Insulin'].quantile(0.75)
```

```
insulinIQR=insulinQ3-insulinQ1
```

```
lowerliminsulin=insulinQ1-1.5*insulinIQR
```

```
upperliminsulin=insulinQ3+1.5*insulinIQR
```

```
insulrem=cancer[(cancer['Insulin']>lowerliminsulin)&(upperliminsulin >  
cancer['Insulin'])]
```

```
sns.boxplot(insulrem['Glucose'])  
glucoseQ1=insulrem['Glucose'].quantile(0.25)  
glucoseQ3=insulrem['Glucose'].quantile(0.75)  
glucoseIQR=glucoseQ3-glucoseQ1  
upperlimglucose=glucoseQ3+1.5*glucoseIQR  
lowerlimglucose=glucoseQ1-1.5*glucoseIQR  
glucosere=insulrem[(insulrem['Glucose'] >  
lowerlimglucose)&(upperlimglucose > insulrem['Glucose'])]
```

```
sns.boxplot(glucoserem['HOMA'])  
homaQ1=glucoserem['HOMA'].quantile(0.25)  
homaQ3=glucoserem['HOMA'].quantile(0.75)  
homaIQR=homaQ3-homaQ1  
upperlimhoma=homaQ3+1.5*homaIQR  
lowerlimhoma=homaQ1-1.5*homaIQR  
homarem=glucoserem[(glucoserem['HOMA'] >  
lowerlimhoma)&(upperlimhoma > glucoserem['HOMA'])]
```

```
sns.boxplot(homarem['Adiponectin'])  
AdiponectinQ1=homarem['Adiponectin'].quantile(0.25)  
AdiponectinQ3=homarem['Adiponectin'].quantile(0.75)  
AdiponectinIQR=AdiponectinQ3-AdiponectinQ1  
upperlimAdiponectin=AdiponectinQ3+1.5*AdiponectinIQR  
lowerlimAdiponectin=AdiponectinQ1-1.5*AdiponectinIQR  
adirem=homarem[(homarem['Adiponectin'] >  
lowerlimAdiponectin)&(upperlimAdiponectin >  
homarem['Adiponectin'])]
```

```
sns.boxplot(adirem['Resistin'])
```



```
sns.boxplot(adirem['Leptin'])
```

```
sns.boxplot(adirem['MCP.1'])
```

```
# create the features from data
```

```
X=mcprem.iloc[:,0:9]
```

```
# create the target variable from data
```

```
Y=mcprem.iloc[:,9]
```

STANDARDIZE

STANDARDIZE TO BRING ALL TO THE SAME SCALE

```
from sklearn.preprocessing import StandardScaler  
  
ss=StandardScaler()  
  
X=ss.fit_transform(X)  
  
X=pd.DataFrame(X)
```

SPLIT DATA

```
from sklearn.model_selection import train_test_split  
  
xtrain,xtest,ytrain,ytest=train_test_split(X,Y,test_size=0.3)
```

BUILD THE CLASSIFIER

BUILD THE MODEL

```
#Finding accuracies on TrainData and Test data with euclidean distance(by default p=2)
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
for x in range(5,10,2):
```

```
    knn=KNeighborsClassifier(n_neighbors=x,metric='minkowski',weights='distance')
```

```
    knn.fit(xtrain,ytrain)
```

```
    train_ypred=knn.predict(xtrain)
```

```
    acc_train_score=accuracy_score(train_ypred,ytrain)
```

```
    test_ypred=knn.predict(xtest)
```

```
    acc_test_score=accuracy_score(test_ypred,ytest)
```

```
    print(f'Accuracy score for train data and test data is {acc_train_score} and {acc_test_score} respectively for {x} neighbours')
```

BUILD THE MODEL WITH EUCLIDEAN DISTANCE

```
knn=KNeighborsClassifier(n_neighbors=7,metric='minkowski',weights='distance')
```

```
knn.fit(xtrain,ytrain)
```

```
trainypred=knn.predict(xtrain)
```

RUN A METRIC REPORT

```
from sklearn.metrics import classification_report  
print(classification_report(trainypred,ytrain))  
  
accuracy_score(trainypred,ytrain)  
  
testypredicted=knn.predict(xtest)  
  
from sklearn.metrics import accuracy_score  
accuracy_score(testypredicted,ytest)
```

	precision	recall	f1-score	support
1	0.92	0.71	0.80	31
2	0.64	0.89	0.74	18
accuracy			0.78	49
macro avg	0.78	0.80	0.77	49
weighted avg	0.82	0.78	0.78	49

PICKLES

```
import pickle

#Save our model as a pickle to a file

pickle.dump(knn, open("my_knn_model.pickle.dat", "wb"))

# delete the existing knn model from the environment

del knn

#Load the pickled object from the file

load_knn=pickle.load(open("my_knn_model.pickle.dat", "rb"))

# Use the loaded model to make predictions

load_knn.predict(xtest)
```

SUMMARY